

Comp 472 Project 1

by Maharaj Teertha Deb (40227747)

✓ Load the data

✓ Importing dataset to devite it between train and test datasets.

```
from keras.datasets import cifar10

(X_train, y_train), (X_test, y_test) = cifar10.load_data()

print("Training data shape:", X_train.shape, y_train.shape)
print("Test data shape:", X_test.shape, y_test.shape)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
 170498071/170498071 — 4s 0us/step
 Training data shape: (50000, 32, 32, 3) (50000, 1)
 Test data shape: (10000, 32, 32, 3) (10000, 1)

✓ Using first 500 training images and the first 100 test images of each class

```
import numpy as np

X_train_filtered = []
y_train_filtered = []
X_test_filtered = []
y_test_filtered = []

# Loop through each class
for class_label in range(10):
    train_indices = np.where(y_train.flatten() == class_label)[0][:500] # 500 images for each class for training purposes
    test_indices = np.where(y_test.flatten() == class_label)[0][:100] # 100 images for each class for test purposes.

    #data to the filtered arrays
    X_train_filtered.append(X_train[train_indices])
    y_train_filtered.append(y_train[train_indices])
    X_test_filtered.append(X_test[test_indices])
    y_test_filtered.append(y_test[test_indices])

# print("X_train: ", X_train_filtered)
# print("y_train: ", y_train_filtered)
# print("X_test: ", X_test_filtered)
# print("y_test: ", y_test_filtered)

# concatenating lists to form final arrays
X_train_filtered = np.concatenate(X_train_filtered)
y_train_filtered = np.concatenate(y_train_filtered)
X_test_filtered = np.concatenate(X_test_filtered)
y_test_filtered = np.concatenate(y_test_filtered)

# the shapes of the filtered datasets
print("training data shape:", X_train_filtered.shape, y_train_filtered.shape)
print("test data shape:", X_test_filtered.shape, y_test_filtered.shape)
```

training data shape: (5000, 32, 32, 3) (5000, 1)
 test data shape: (1000, 32, 32, 3) (1000, 1)

✓ Here we do:

1. Resizing and

2. Normalizing

** imageNet means and std were taken from: [stack-overflow](#)

```
import torch
import torchvision.transforms as transforms

transform = transforms.Compose([
    transforms.ToPILImage(), # Convert to PIL image
    transforms.Resize((224, 224)), # Resize to 224x224
    transforms.ToTensor(), # Convert to tensor
    transforms.Normalize(mean=[0.485, 0.456, 0.406], # ImageNet means
                        std=[0.229, 0.224, 0.225]) # ImageNet stds
])

# now applying the transformation to the training and test sets
X_train_normalized = torch.stack([transform(img) for img in X_train_filtered])
X_test_normalized = torch.stack([transform(img) for img in X_test_filtered])

print("Training data shape after transformation:", X_train_normalized.shape)
print("Test data shape after transformation:", X_test_normalized.shape)
```

```
↗ Training data shape after transformation: torch.Size([5000, 3, 224, 224])
Test data shape after transformation: torch.Size([1000, 3, 224, 224])
```

✓ Now we can extract features using ResNet-18

device will be selected based on GPU and CPU. original code: [Pytorch website](#)

how to use pretrained resnet and evalute it, was taken from [Stack-Overflow](#)

```
from torch.utils.data import DataLoader, TensorDataset
from torchvision.models import resnet18

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu') # taken from pytorch

resnet18 = resnet18(pretrained=True).to(device)
resnet18.eval() # stackoverflow code

batch_size = 32 # with higher batch size, my colab environment crashes...
train_dataset = TensorDataset(X_train_normalized)
test_dataset = TensorDataset(X_test_normalized)
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

# Function to extract features in batches
def extract_features_in_batches(model, loader):
    all_features = []
    for batch in loader:
        data = batch[0].to(device) # Move data to the same device as the model
        with torch.no_grad(): # No gradient computation for inference
            features = model(data).squeeze() # Extract features
        all_features.append(features.cpu()) # Move to CPU to free GPU memory
    return torch.cat(all_features)

# Now we extract features for training and test sets
X_train_features = extract_features_in_batches(resnet18, train_loader)
X_test_features = extract_features_in_batches(resnet18, test_loader)

print("Feature extraction complete with batch processing.")
print("Training features shape:", X_train_features.shape)
print("Test features shape:", X_test_features.shape)
```

```
↗ /usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|██████████| 44.7M/44.7M [00:00<00:00, 63.4MB/s]
```

✓ Final task for Dataset overview, PCA Dimensionality Reduction

```
from sklearn.decomposition import PCA

pca = PCA(n_components=50)

# Applying PCA on the training and test feature sets
X_train_pca = pca.fit_transform(X_train_features.cpu().numpy())
X_test_pca = pca.transform(X_test_features.cpu().numpy())

print("Training features shape after PCA:", X_train_pca.shape)
print("Test features shape after PCA:", X_test_pca.shape)
```

✓ Naive Bayes:

✓ Custom Gaussian Naive Bayes Implementation

idea was taken from:

1. [StackOverflow](#)
2. [Scikit-learn's naive_bayes class](#)

** Also Google GEMINI Suggested some codes(like co-pilot) while writing it on colab environment.

```
import numpy as np
```

```
class GaussianNaiveBayes:
    def __init__(self) -> None:
        self.means = None
        self.variances = None
        self.priors = None

    def fit(self, X, y) -> None:
        self.classes = np.unique(y)
        n_features = X.shape[1]
        n_classes = len(self.classes)

        self.means = np.zeros((n_classes, n_features))
        self.variances = np.zeros((n_classes, n_features))
        self.priors = np.zeros(n_classes)

        for idx, cls in enumerate(self.classes):
            X_cls = X[y == cls]
            self.means[idx, :] = X_cls.mean(axis=0)
            self.variances[idx, :] = X_cls.var(axis=0) + 1e-9 # Added small value to avoid division by zero ** This idea is from GOOGLE GEM
            self.priors[idx] = X_cls.shape[0] / X.shape[0]

    def predict(self, X) -> np.ndarray:
        posteriors = np.zeros((X.shape[0], len(self.classes)))

        for idx, cls in enumerate(self.classes):
            prior = np.log(self.priors[idx])
            likelihood = -0.5 * np.sum(np.log(2 * np.pi * self.variances[idx, :]))
            likelihood -= 0.5 * np.sum(((X - self.means[idx, :]) ** 2) / self.variances[idx, :], axis=1)
            posteriors[:, idx] = prior + likelihood

        return self.classes[np.argmax(posteriors, axis=1)]

gnb_custom = GaussianNaiveBayes()
gnb_custom.fit(X_train_pca, y_train_filtered.flatten()) # Flatten y_train for compatibility

y_pred_custom = gnb_custom.predict(X_test_pca)
print("Custom Gaussian Naive Bayes completed.")
```


➡ Custom Gaussian Naive Bayes completed.

✓ Gaussian Naive Bayes Using Scikit-Learn

```
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

gnb_sklearn = GaussianNB()
gnb_sklearn.fit(X_train_pca, y_train_filtered.flatten())

y_pred_sklearn = gnb_sklearn.predict(X_test_pca)
print("Scikit-learn Gaussian Naive Bayes completed.")
```

 Scikit-learn Gaussian Naive Bayes completed.

✓ Evaluate both of these models on the test set.

```
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay

accuracy_custom = accuracy_score(y_test_filtered.flatten(), y_pred_custom)
accuracy_sklearn = accuracy_score(y_test_filtered.flatten(), y_pred_sklearn)
```

```
print("\nCustom GNB Classification Report:")
print(classification_report(y_test_filtered.flatten(), y_pred_custom))
```

```
print("\nScikit-learn GNB Classification Report:")
print(classification_report(y_test_filtered.flatten(), y_pred_sklearn))
```

```
fig, axes = plt.subplots(1, 2, figsize=(14, 6)) # confusion matrix
```

```
ConfusionMatrixDisplay.from_predictions(
    y_test_filtered.flatten(),
    y_pred_custom,
    ax=axes[0],
    cmap="Blues",
    colorbar=False
)
axes[0].set_title("Custom GNB")
```

```
ConfusionMatrixDisplay.from_predictions(
    y_test_filtered.flatten(),
    y_pred_sklearn,
    ax=axes[1],
    cmap="Greens",
    colorbar=False
)
axes[1].set_title("Scikit-learn GNB")
```

```
plt.tight_layout()
plt.show()
```

```
models = ['Custom GNB', 'Scikit-learn GNB']
accuracies = [accuracy_custom, accuracy_sklearn]
```

```
plt.figure(figsize=(8, 5))
plt.bar(models, accuracies, color=['blue', 'green'])
plt.title('Model Accuracy Comparison')
plt.ylabel('Accuracy')
plt.ylim(0, 1)
plt.show()
```

```
# now i save the naive_bayes_model::
import pickle
with open('naive_bayes_model.pkl', 'wb') as f:
    pickle.dump(gnb_custom, f)
    print("Model saved to naive_bayes_model.pkl")
from google.colab import files
files.download('naive_bayes_model.pkl')
```


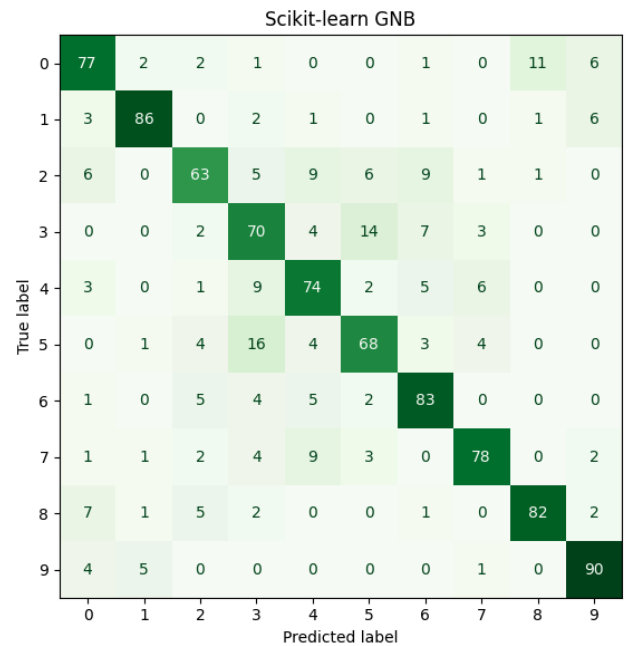




Custom GNB Classification Report:

	precision	recall	f1-score	support
0	0.75	0.77	0.76	100
1	0.90	0.86	0.88	100
2	0.75	0.63	0.68	100
3	0.62	0.70	0.66	100
4	0.70	0.74	0.72	100
5	0.72	0.68	0.70	100
6	0.75	0.83	0.79	100
7	0.84	0.78	0.81	100
8	0.86	0.82	0.84	100
9	0.85	0.90	0.87	100
accuracy			0.77	1000
macro avg	0.77	0.77	0.77	1000
weighted avg	0.77	0.77	0.77	1000

Scikit-learn GNB Classification Report:

	precision	recall	f1-score	support
0	0.75	0.77	0.76	100
1	0.90	0.86	0.88	100
2	0.75	0.63	0.68	100
3	0.62	0.70	0.66	100
4	0.70	0.74	0.72	100
5	0.72	0.68	0.70	100
6	0.75	0.83	0.79	100
7	0.84	0.78	0.81	100
8	0.86	0.82	0.84	100
9	0.85	0.90	0.87	100
accuracy			0.77	1000
macro avg	0.77	0.77	0.77	1000
weighted avg	0.77	0.77	0.77	1000





✓ Section 4: Decision Tree.

✓ Custom Decision Tree Classifier (Gini Index, Max Depth = 50)

the idea was taken from: [Ander Fernandez's blog](#) and [Medium.com](#)

```
import numpy as np
```

```
class DecisionTreeClassifier:
    def __init__(self, max_depth=50):
        self.max_depth = max_depth
        self.tree = None

    def gini(self, y):
        m = len(y)
        if m == 0:
            return 0
        _, counts = np.unique(y, return_counts=True)
        probabilities = counts / m
        return 1 - np.sum(probabilities ** 2)

    def split(self, X, y, feature, threshold):
        left_indices = X[:, feature] < threshold
        right_indices = ~left_indices
        return X[left_indices], y[left_indices], X[right_indices], y[right_indices]

    def best_split(self, X, y):
        m, n = X.shape
        if m <= 1:
            return None, None
        parent_gini = self.gini(y)
        best_gain = 0
        best_split_feature, best_split_threshold = None, None
        for feature in range(n):
            thresholds = np.unique(X[:, feature])
            for threshold in thresholds:
                X_left, y_left, X_right, y_right = self.split(X, y, feature, threshold)
                if len(y_left) == 0 or len(y_right) == 0:
                    continue
                n_left, n_right = len(y_left), len(y_right)
                child_gini = (n_left / m) * self.gini(y_left) + (n_right / m) * self.gini(y_right)
                gain = parent_gini - child_gini
                if gain > best_gain:
                    best_gain = gain
                    best_split_feature, best_split_threshold = feature, threshold
        return best_split_feature, best_split_threshold

    def build_tree(self, X, y, depth=0):
        if depth >= self.max_depth or len(np.unique(y)) == 1:
            return np.argmax(np.bincount(y))
        feature, threshold = self.best_split(X, y)
        if feature is None:
            return np.argmax(np.bincount(y))
        X_left, y_left, X_right, y_right = self.split(X, y, feature, threshold)
        return {
            'feature': feature,
            'threshold': threshold,
```

```

        'left': self.build_tree(X_left, y_left, depth + 1),
        'right': self.build_tree(X_right, y_right, depth + 1),
    }

    def fit(self, X, y):
        self.tree = self.build_tree(X, y)

    def predict_sample(self, sample, tree):
        if isinstance(tree, dict):
            if sample[tree['feature']] < tree['threshold']:
                return self.predict_sample(sample, tree['left'])
            else:
                return self.predict_sample(sample, tree['right'])
        else:
            return tree

    def predict(self, X):
        return np.array([self.predict_sample(sample, self.tree) for sample in X])

```

✓ Train the tree:

```

custom_tree = DecisionTreeClassifier(max_depth=50)
custom_tree.fit(X_train_pca, y_train_filtered.flatten())

```

✓ Save the model:

**** code suggested by GOOGLE GEMINI. propt was: I want to download my custom decision tree.**

```


import cloudpickle

# Save the custom tree model using cloudpickle
with open('dtc_cloudpickle.pkl', 'wb') as f:
    cloudpickle.dump(custom_tree, f)

print("Model saved to dtc_cloudpickle.pkl")

# Download the model file
from google.colab import files
files.download('dtc_cloudpickle.pkl')

```

 Model saved to dtc_cloudpickle.pkl

✓ Predict on Test Set:

```

y_pred_custom_tree = custom_tree.predict(X_test_pca)

```

✓ Experimenting with Different Depths

```

depths = [5, 10, 15, 20, 35, 50] # We will experiment with 6 depths from 5 to 50
accuracies_for_custom_tree = []

```

```

for depth in depths:
    custom_tree = DecisionTreeClassifier(max_depth=depth)
    custom_tree.fit(X_train_pca, y_train_filtered.flatten())

    y_pred = custom_tree.predict(X_test_pca)
    #print(y_pred)

    accuracy = np.mean(y_pred == y_test_filtered.flatten())

    accuracies_for_custom_tree.append(accuracy)
    print(f"Tree Depth: {depth}, Accuracy: {accuracy:.4f}")

```


```

import matplotlib.pyplot as plt

```

```
plt.figure(figsize=(10, 6))
plt.plot(depths, accuracies_for_custome_tree, marker='o', label='Test Accuracy')
plt.title("Effect of Tree Depth on Model Accuracy")
plt.xlabel("Tree Depth")
plt.ylabel("Accuracy")
plt.grid(True)
plt.legend()
plt.show()
```

```
Tree Depth: 5, Accuracy: 0.5850
Tree Depth: 10, Accuracy: 0.6030
Tree Depth: 15, Accuracy: 0.5800
Tree Depth: 20, Accuracy: 0.5780
Tree Depth: 35, Accuracy: 0.5780
Tree Depth: 50, Accuracy: 0.5780
```



✓ Scikit-learn's DecisionTreeClassifier

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

accuracies_for_sklearn_tree = []

for depth in depths:
    sk_tree = DecisionTreeClassifier(criterion="gini", max_depth=depth, random_state=42)
    sk_tree.fit(X_train_pca, y_train_filtered.flatten())

    y_pred_sklearn_tree = sk_tree.predict(X_test_pca)

    accuracy = accuracy_score(y_test_filtered.flatten(), y_pred_sklearn_tree)
    accuracies_for_sklearn_tree.append(accuracy)
    print(f"Tree Depth: {depth}, Accuracy: {accuracy:.4f}")


plt.figure(figsize=(10, 6))
plt.plot(depths, accuracies_for_sklearn_tree, marker='o', label='Test Accuracy')
plt.title("Effect of Tree Depth on Model Accuracy (Scikit-Learn)")
plt.xlabel("Tree Depth")
plt.ylabel("Accuracy")
plt.grid(True)
plt.legend()
plt.show()
```



```

↩ Tree Depth: 5, Accuracy: 0.5860
Tree Depth: 10, Accuracy: 0.6100
Tree Depth: 15, Accuracy: 0.6050
Tree Depth: 20, Accuracy: 0.6000
Tree Depth: 35, Accuracy: 0.6000
Tree Depth: 50, Accuracy: 0.6000

```



✓ Evaluation and Graphical Comparison

```

from sklearn.metrics import classification_report, ConfusionMatrixDisplay

print("Custom Decision Tree Classification Report:")
print(classification_report(y_test_filtered.flatten(), y_pred_custom_tree))

print("\nScikit-Learn Decision Tree Classification Report:")
print(classification_report(y_test_filtered.flatten(), y_pred_sklearn_tree))

fig, axes = plt.subplots(1, 2, figsize=(14, 6)) # for plotting confusion matrix

ConfusionMatrixDisplay.from_predictions(
    y_test_filtered.flatten(),
    y_pred_custom_tree,
    ax=axes[0],
    cmap="Blues",
    colorbar=False
)
axes[0].set_title("Custom Decision Tree")

# Scikit-Learn Decision Tree confusion matrix
ConfusionMatrixDisplay.from_predictions(
    y_test_filtered.flatten(),
    y_pred_sklearn_tree,
    ax=axes[1],
    cmap="Greens",
    colorbar=False
)
axes[1].set_title("Scikit-Learn Decision Tree")

plt.tight_layout()
plt.show()



```

Custom Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.59	0.55	0.57	100
1	0.73	0.72	0.72	100
2	0.43	0.40	0.42	100
3	0.42	0.43	0.42	100
4	0.48	0.53	0.50	100
5	0.56	0.61	0.59	100
6	0.67	0.69	0.68	100
7	0.49	0.48	0.48	100
8	0.63	0.63	0.63	100
9	0.80	0.74	0.77	100
accuracy			0.58	1000
macro avg	0.58	0.58	0.58	1000
weighted avg	0.58	0.58	0.58	1000

Scikit-Learn Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.57	0.60	0.59	100
1	0.78	0.77	0.77	100
2	0.47	0.40	0.43	100
3	0.48	0.48	0.48	100
4	0.50	0.51	0.50	100
5	0.58	0.62	0.60	100
6	0.74	0.67	0.71	100
7	0.48	0.49	0.48	100
8	0.64	0.69	0.66	100
9	0.75	0.77	0.76	100
accuracy			0.60	1000
macro avg	0.60	0.60	0.60	1000
weighted avg	0.60	0.60	0.60	1000



MLP

Implementation of the specified three-layer MLP using PyTorch:

Help taken from: [pytorch discussion](#) and [medium.com](#)

I just ignored the layers part. as we need to implement three-layerd MLP.

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(50, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.BatchNorm1d(512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        return self.model(x)

batch_size = 32

train_dataset = TensorDataset(torch.tensor(X_train_pca, dtype=torch.float32),
                               torch.tensor(y_train_filtered.flatten(), dtype=torch.long))
test_dataset = TensorDataset(torch.tensor(X_test_pca, dtype=torch.float32),
                              torch.tensor(y_test_filtered.flatten(), dtype=torch.long))

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

# device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

mlp_model = MLP().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(mlp_model.parameters(), lr=0.01, momentum=0.9)

num_epochs = 100
for epoch in range(num_epochs):
    mlp_model.train()
    running_loss = 0.0
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)

        # Forward pass
        outputs = mlp_model(inputs)
        loss = criterion(outputs, labels)

        # Backward pass and optimization
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(train_loader):.4f}")

mlp_model.eval()
correct = 0
total = 0

with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = mlp_model(inputs)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = correct / total
print(f"Test Accuracy: {accuracy:.4f}")

#now we save the model:
## Code suggested by Google GEMINI.
```

```

torch.save(mlp_model.state_dict(), 'mlp_model.pth')
print("Model saved to mlp_model.pth")
# Download the model file
from google.colab import files
files.download('mlp_model.pth')

```

```

Epoch [44/100], Loss: 0.0040
Epoch [45/100], Loss: 0.0051
Epoch [46/100], Loss: 0.0099
Epoch [47/100], Loss: 0.0074
Epoch [48/100], Loss: 0.0063
Epoch [49/100], Loss: 0.0185
Epoch [50/100], Loss: 0.0187
Epoch [51/100], Loss: 0.0241
Epoch [52/100], Loss: 0.0331
Epoch [53/100], Loss: 0.0168
Epoch [54/100], Loss: 0.0379
Epoch [55/100], Loss: 0.0282
Epoch [56/100], Loss: 0.0279
Epoch [57/100], Loss: 0.0195
Epoch [58/100], Loss: 0.0117
Epoch [59/100], Loss: 0.0112
Epoch [60/100], Loss: 0.0066
Epoch [61/100], Loss: 0.0037
Epoch [62/100], Loss: 0.0031
Epoch [63/100], Loss: 0.0065
Epoch [64/100], Loss: 0.0033
Epoch [65/100], Loss: 0.0029
Epoch [66/100], Loss: 0.0028
Epoch [67/100], Loss: 0.0020
Epoch [68/100], Loss: 0.0030
Epoch [69/100], Loss: 0.0020
Epoch [70/100], Loss: 0.0012
Epoch [71/100], Loss: 0.0014
Epoch [72/100], Loss: 0.0011
Epoch [73/100], Loss: 0.0028
Epoch [74/100], Loss: 0.0190
Epoch [75/100], Loss: 0.0384
Epoch [76/100], Loss: 0.0116
Epoch [77/100], Loss: 0.0041
Epoch [78/100], Loss: 0.0034
Epoch [79/100], Loss: 0.0024
Epoch [80/100], Loss: 0.0031
Epoch [81/100], Loss: 0.0021
Epoch [82/100], Loss: 0.0028
Epoch [83/100], Loss: 0.0166
Epoch [84/100], Loss: 0.0468
Epoch [85/100], Loss: 0.0418
Epoch [86/100], Loss: 0.0263
Epoch [87/100], Loss: 0.0159
Epoch [88/100], Loss: 0.0096
Epoch [89/100], Loss: 0.0045
Epoch [90/100], Loss: 0.0044
Epoch [91/100], Loss: 0.0119
Epoch [92/100], Loss: 0.0043
Epoch [93/100], Loss: 0.0028
Epoch [94/100], Loss: 0.0021
Epoch [95/100], Loss: 0.0017
Epoch [96/100], Loss: 0.0018
Epoch [97/100], Loss: 0.0028
Epoch [98/100], Loss: 0.0026
Epoch [99/100], Loss: 0.0018
Epoch [100/100], Loss: 0.0023
Test Accuracy: 0.8070

```

✓ experiment by varying the depth of the MLP and document the influence on performance:

```

class MLPVariant(nn.Module):
    def __init__(self, depth):
        super(MLPVariant, self).__init__()
        layers = []

        layers.append(nn.Linear(50, 512))
        layers.append(nn.ReLU())

        for _ in range(depth - 2): # -2 to account for input and output layers
            layers.append(nn.Linear(512, 512))

```

```

        layers.append(nn.BatchNorm1d(512))
        layers.append(nn.ReLU())

    layers.append(nn.Linear(512, 10))

    self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)

depths = [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20]
accuracies = []

for depth in depths:
    model = MLPVariant(depth).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

    num_epochs = 10
    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            running_loss += loss.item()

        model.eval()
        correct = 0
        total = 0
        with torch.no_grad():
            for inputs, labels in test_loader:
                inputs, labels = inputs.to(device), labels.to(device)
                outputs = model(inputs)
                _, predicted = torch.max(outputs, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

        accuracy = correct / total
        print(f"Epoch [{epoch+1}/{num_epochs}], Depth: {depth}, Loss: {running_loss/len(train_loader):.4f}, Accuracy: {accuracy:.4f}")

    accuracies.append(accuracy)
    print(f"Final Depth: {depth}, Test Accuracy: {accuracy:.4f}")

import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.plot(depths, accuracies, marker='o', label="Test Accuracy")
plt.title("Effect of MLP Depth on Test Accuracy")
plt.xlabel("Number of Layers")
plt.ylabel("Accuracy")
plt.grid(True)
plt.legend()
plt.show()

```

```


Epoch [1/10], Depth: 2, Loss: 0.8791, Accuracy: 0.7740
Epoch [2/10], Depth: 2, Loss: 0.5201, Accuracy: 0.7740
Epoch [3/10], Depth: 2, Loss: 0.4757, Accuracy: 0.8060
Epoch [4/10], Depth: 2, Loss: 0.4122, Accuracy: 0.8180
Epoch [5/10], Depth: 2, Loss: 0.3724, Accuracy: 0.8050
Epoch [6/10], Depth: 2, Loss: 0.3435, Accuracy: 0.7980
Epoch [7/10], Depth: 2, Loss: 0.3285, Accuracy: 0.8240
Epoch [8/10], Depth: 2, Loss: 0.2791, Accuracy: 0.8160
Epoch [9/10], Depth: 2, Loss: 0.2768, Accuracy: 0.8160
Epoch [10/10], Depth: 2, Loss: 0.2468, Accuracy: 0.8170
Final Depth: 2, Test Accuracy: 0.8170
Epoch [1/10], Depth: 3, Loss: 1.1277, Accuracy: 0.7630
Epoch [2/10], Depth: 3, Loss: 0.6654, Accuracy: 0.7840
Epoch [3/10], Depth: 3, Loss: 0.5618, Accuracy: 0.8000
Epoch [4/10], Depth: 3, Loss: 0.5083, Accuracy: 0.8140
Epoch [5/10], Depth: 3, Loss: 0.4646, Accuracy: 0.8180
Epoch [6/10], Depth: 3, Loss: 0.4313, Accuracy: 0.8170
Epoch [7/10], Depth: 3, Loss: 0.4111, Accuracy: 0.8290
Epoch [8/10], Depth: 3, Loss: 0.3851, Accuracy: 0.8240
Epoch [9/10], Depth: 3, Loss: 0.3543, Accuracy: 0.8340
Epoch [10/10], Depth: 3, Loss: 0.3375, Accuracy: 0.8260
Final Depth: 3, Test Accuracy: 0.8260
Epoch [1/10], Depth: 4, Loss: 1.0506, Accuracy: 0.7750
Epoch [2/10], Depth: 4, Loss: 0.6120, Accuracy: 0.7990
Epoch [3/10], Depth: 4, Loss: 0.5097, Accuracy: 0.8050
Epoch [4/10], Depth: 4, Loss: 0.4436, Accuracy: 0.8230
Epoch [5/10], Depth: 4, Loss: 0.3906, Accuracy: 0.8140
Epoch [6/10], Depth: 4, Loss: 0.3548, Accuracy: 0.8160
Epoch [7/10], Depth: 4, Loss: 0.3017, Accuracy: 0.8160
Epoch [8/10], Depth: 4, Loss: 0.2657, Accuracy: 0.8170
Epoch [9/10], Depth: 4, Loss: 0.2329, Accuracy: 0.8270
Epoch [10/10], Depth: 4, Loss: 0.2087, Accuracy: 0.8250
Final Depth: 4, Test Accuracy: 0.8250
Epoch [1/10], Depth: 5, Loss: 1.0724, Accuracy: 0.7640
Epoch [2/10], Depth: 5, Loss: 0.5944, Accuracy: 0.7970
Epoch [3/10], Depth: 5, Loss: 0.4743, Accuracy: 0.7950
Epoch [4/10], Depth: 5, Loss: 0.3942, Accuracy: 0.8010
Epoch [5/10], Depth: 5, Loss: 0.3372, Accuracy: 0.8140
Epoch [6/10], Depth: 5, Loss: 0.2822, Accuracy: 0.8140
Epoch [7/10], Depth: 5, Loss: 0.2283, Accuracy: 0.8180
Epoch [8/10], Depth: 5, Loss: 0.1916, Accuracy: 0.8160
Epoch [9/10], Depth: 5, Loss: 0.1601, Accuracy: 0.8110
Epoch [10/10], Depth: 5, Loss: 0.1341, Accuracy: 0.8030
Final Depth: 5, Test Accuracy: 0.8030
Epoch [1/10], Depth: 6, Loss: 1.0474, Accuracy: 0.7830
Epoch [2/10], Depth: 6, Loss: 0.5766, Accuracy: 0.8040
Epoch [3/10], Depth: 6, Loss: 0.4514, Accuracy: 0.8270
Epoch [4/10], Depth: 6, Loss: 0.3532, Accuracy: 0.8260
Epoch [5/10], Depth: 6, Loss: 0.2859, Accuracy: 0.8150
Epoch [6/10], Depth: 6, Loss: 0.2239, Accuracy: 0.8050
Epoch [7/10], Depth: 6, Loss: 0.1905, Accuracy: 0.8170
Epoch [8/10], Depth: 6, Loss: 0.1373, Accuracy: 0.8160
Epoch [9/10], Depth: 6, Loss: 0.1131, Accuracy: 0.8240
Epoch [10/10], Depth: 6, Loss: 0.0866, Accuracy: 0.8180
Final Depth: 6, Test Accuracy: 0.8180
Epoch [1/10], Depth: 7, Loss: 1.0998, Accuracy: 0.7700
Epoch [2/10], Depth: 7, Loss: 0.5950, Accuracy: 0.8020
Epoch [3/10], Depth: 7, Loss: 0.4404, Accuracy: 0.7970
Epoch [4/10], Depth: 7, Loss: 0.3441, Accuracy: 0.8010
Epoch [5/10], Depth: 7, Loss: 0.2635, Accuracy: 0.8070
Epoch [6/10], Depth: 7, Loss: 0.2077, Accuracy: 0.7970
Epoch [7/10], Depth: 7, Loss: 0.1668, Accuracy: 0.8020
Epoch [8/10], Depth: 7, Loss: 0.1409, Accuracy: 0.8060
Epoch [9/10], Depth: 7, Loss: 0.1035, Accuracy: 0.8090
Epoch [10/10], Depth: 7, Loss: 0.0936, Accuracy: 0.7920
Final Depth: 7, Test Accuracy: 0.7920
Epoch [1/10], Depth: 8, Loss: 1.1075, Accuracy: 0.7770
Epoch [2/10], Depth: 8, Loss: 0.5813, Accuracy: 0.8060
Epoch [3/10], Depth: 8, Loss: 0.4488, Accuracy: 0.7930
Epoch [4/10], Depth: 8, Loss: 0.3531, Accuracy: 0.7980
Epoch [5/10], Depth: 8, Loss: 0.2727, Accuracy: 0.8170
Epoch [6/10], Depth: 8, Loss: 0.2053, Accuracy: 0.8030
Epoch [7/10], Depth: 8, Loss: 0.1717, Accuracy: 0.8100
Epoch [8/10], Depth: 8, Loss: 0.1378, Accuracy: 0.7910
Epoch [9/10], Depth: 8, Loss: 0.1021, Accuracy: 0.8040
Epoch [10/10], Depth: 8, Loss: 0.0955, Accuracy: 0.7780
Final Depth: 8, Test Accuracy: 0.7780
Epoch [1/10], Depth: 9, Loss: 1.1652, Accuracy: 0.7780
Epoch [2/10], Depth: 9, Loss: 0.6125, Accuracy: 0.7930
Epoch [3/10], Depth: 9, Loss: 0.4748, Accuracy: 0.7940
Epoch [4/10], Depth: 9, Loss: 0.3843, Accuracy: 0.7920
Epoch [5/10], Depth: 9, Loss: 0.3149, Accuracy: 0.7940
Epoch [6/10], Depth: 9, Loss: 0.2581, Accuracy: 0.7980
Epoch [7/10], Depth: 9, Loss: 0.2038, Accuracy: 0.7790
Epoch [8/10], Depth: 9, Loss: 0.1555, Accuracy: 0.7790
Epoch [9/10], Depth: 9, Loss: 0.1072, Accuracy: 0.7790
Epoch [10/10], Depth: 9, Loss: 0.0955, Accuracy: 0.7790
Final Depth: 9, Test Accuracy: 0.7790

```

```
Epoch [8/10], Depth: 9, Loss: 0.1580, Accuracy: 0.7960
Epoch [9/10], Depth: 9, Loss: 0.1297, Accuracy: 0.7990
Epoch [10/10], Depth: 9, Loss: 0.1210, Accuracy: 0.7920
Final Depth: 9, Test Accuracy: 0.7920
Epoch [1/10], Depth: 10, Loss: 1.1944, Accuracy: 0.7610
Epoch [2/10], Depth: 10, Loss: 0.6291, Accuracy: 0.7810
Epoch [3/10], Depth: 10, Loss: 0.5185, Accuracy: 0.7940
Epoch [4/10], Depth: 10, Loss: 0.4026, Accuracy: 0.7940
Epoch [5/10], Depth: 10, Loss: 0.3081, Accuracy: 0.8040
Epoch [6/10], Depth: 10, Loss: 0.2585, Accuracy: 0.8020
Epoch [7/10], Depth: 10, Loss: 0.2237, Accuracy: 0.7890
Epoch [8/10], Depth: 10, Loss: 0.1936, Accuracy: 0.7930
Epoch [9/10], Depth: 10, Loss: 0.1809, Accuracy: 0.7860
Epoch [10/10], Depth: 10, Loss: 0.1597, Accuracy: 0.7920
Final Depth: 10, Test Accuracy: 0.7920
Epoch [1/10], Depth: 11, Loss: 1.2556, Accuracy: 0.7610
Epoch [2/10], Depth: 11, Loss: 0.6739, Accuracy: 0.7750
Epoch [3/10], Depth: 11, Loss: 0.5504, Accuracy: 0.7870
Epoch [4/10], Depth: 11, Loss: 0.4378, Accuracy: 0.7890
Epoch [5/10], Depth: 11, Loss: 0.3626, Accuracy: 0.7900
Epoch [6/10], Depth: 11, Loss: 0.3331, Accuracy: 0.7710
Epoch [7/10], Depth: 11, Loss: 0.2650, Accuracy: 0.7730
Epoch [8/10], Depth: 11, Loss: 0.2382, Accuracy: 0.7830
Epoch [9/10], Depth: 11, Loss: 0.2303, Accuracy: 0.7730
Epoch [10/10], Depth: 11, Loss: 0.1731, Accuracy: 0.7770
Final Depth: 11, Test Accuracy: 0.7770
Epoch [1/10], Depth: 12, Loss: 1.2609, Accuracy: 0.7580
Epoch [2/10], Depth: 12, Loss: 0.7062, Accuracy: 0.7960
Epoch [3/10], Depth: 12, Loss: 0.5665, Accuracy: 0.7820
Epoch [4/10], Depth: 12, Loss: 0.4741, Accuracy: 0.8050
Epoch [5/10], Depth: 12, Loss: 0.3988, Accuracy: 0.7990
Epoch [6/10], Depth: 12, Loss: 0.3563, Accuracy: 0.7930
Epoch [7/10], Depth: 12, Loss: 0.3437, Accuracy: 0.7820
Epoch [8/10], Depth: 12, Loss: 0.2991, Accuracy: 0.7870
Epoch [9/10], Depth: 12, Loss: 0.2615, Accuracy: 0.7830
Epoch [10/10], Depth: 12, Loss: 0.2095, Accuracy: 0.7830
Final Depth: 12, Test Accuracy: 0.7830
Epoch [1/10], Depth: 13, Loss: 1.3126, Accuracy: 0.7460
Epoch [2/10], Depth: 13, Loss: 0.7417, Accuracy: 0.7930
Epoch [3/10], Depth: 13, Loss: 0.6067, Accuracy: 0.7850
Epoch [4/10], Depth: 13, Loss: 0.5189, Accuracy: 0.7970
Epoch [5/10], Depth: 13, Loss: 0.4776, Accuracy: 0.7820
Epoch [6/10], Depth: 13, Loss: 0.3934, Accuracy: 0.7850
Epoch [7/10], Depth: 13, Loss: 0.3609, Accuracy: 0.7960
Epoch [8/10], Depth: 13, Loss: 0.2942, Accuracy: 0.7800
Epoch [9/10], Depth: 13, Loss: 0.2840, Accuracy: 0.7930
Epoch [10/10], Depth: 13, Loss: 0.2312, Accuracy: 0.7930
Final Depth: 13, Test Accuracy: 0.7930
Epoch [1/10], Depth: 14, Loss: 1.4029, Accuracy: 0.7370
Epoch [2/10], Depth: 14, Loss: 0.8240, Accuracy: 0.7380
Epoch [3/10], Depth: 14, Loss: 0.6595, Accuracy: 0.7810
Epoch [4/10], Depth: 14, Loss: 0.5575, Accuracy: 0.7820
Epoch [5/10], Depth: 14, Loss: 0.4876, Accuracy: 0.7790
Epoch [6/10], Depth: 14, Loss: 0.4465, Accuracy: 0.7680
Epoch [7/10], Depth: 14, Loss: 0.3720, Accuracy: 0.7830
Epoch [8/10], Depth: 14, Loss: 0.3673, Accuracy: 0.7860
Epoch [9/10], Depth: 14, Loss: 0.3379, Accuracy: 0.7720
Epoch [10/10], Depth: 14, Loss: 0.2710, Accuracy: 0.7900
Final Depth: 14, Test Accuracy: 0.7900
Epoch [1/10], Depth: 15, Loss: 1.4646, Accuracy: 0.6920
Epoch [2/10], Depth: 15, Loss: 0.8482, Accuracy: 0.7560
Epoch [3/10], Depth: 15, Loss: 0.7074, Accuracy: 0.7530
Epoch [4/10], Depth: 15, Loss: 0.5971, Accuracy: 0.7750
Epoch [5/10], Depth: 15, Loss: 0.5412, Accuracy: 0.7670
Epoch [6/10], Depth: 15, Loss: 0.4928, Accuracy: 0.7850
Epoch [7/10], Depth: 15, Loss: 0.4288, Accuracy: 0.7870
Epoch [8/10], Depth: 15, Loss: 0.3843, Accuracy: 0.7740
Epoch [9/10], Depth: 15, Loss: 0.3346, Accuracy: 0.7850
Epoch [10/10], Depth: 15, Loss: 0.3120, Accuracy: 0.7900
Final Depth: 15, Test Accuracy: 0.7900
Epoch [1/10], Depth: 16, Loss: 1.4949, Accuracy: 0.7010
Epoch [2/10], Depth: 16, Loss: 0.9126, Accuracy: 0.7460
Epoch [3/10], Depth: 16, Loss: 0.7441, Accuracy: 0.7650
Epoch [4/10], Depth: 16, Loss: 0.6435, Accuracy: 0.7750
Epoch [5/10], Depth: 16, Loss: 0.5903, Accuracy: 0.7750
Epoch [6/10], Depth: 16, Loss: 0.5088, Accuracy: 0.7870
Epoch [7/10], Depth: 16, Loss: 0.4231, Accuracy: 0.7720
Epoch [8/10], Depth: 16, Loss: 0.4119, Accuracy: 0.7950
Epoch [9/10], Depth: 16, Loss: 0.3846, Accuracy: 0.7600
Epoch [10/10], Depth: 16, Loss: 0.3504, Accuracy: 0.7960
Final Depth: 16, Test Accuracy: 0.7960
Epoch [1/10], Depth: 17, Loss: 1.5057, Accuracy: 0.6760
Epoch [2/10], Depth: 17, Loss: 0.9184, Accuracy: 0.7430
Epoch [3/10], Depth: 17, Loss: 0.7862, Accuracy: 0.7630
Epoch [4/10], Depth: 17, Loss: 0.7013, Accuracy: 0.7600
```

```
Epoch [5/10], Depth: 17, Loss: 0.6005, Accuracy: 0.7810
Epoch [6/10], Depth: 17, Loss: 0.5717, Accuracy: 0.7830
Epoch [7/10], Depth: 17, Loss: 0.5315, Accuracy: 0.7850
Epoch [8/10], Depth: 17, Loss: 0.4548, Accuracy: 0.7900
Epoch [9/10], Depth: 17, Loss: 0.4219, Accuracy: 0.7790
Epoch [10/10], Depth: 17, Loss: 0.3731, Accuracy: 0.7750
Final Depth: 17, Test Accuracy: 0.7750
Epoch [1/10], Depth: 18, Loss: 1.5274, Accuracy: 0.6720
Epoch [2/10], Depth: 18, Loss: 0.9520, Accuracy: 0.7440
Epoch [3/10], Depth: 18, Loss: 0.8378, Accuracy: 0.7480
Epoch [4/10], Depth: 18, Loss: 0.7235, Accuracy: 0.7720
Epoch [5/10], Depth: 18, Loss: 0.6389, Accuracy: 0.7740
Epoch [6/10], Depth: 18, Loss: 0.5854, Accuracy: 0.7860
Epoch [7/10], Depth: 18, Loss: 0.5511, Accuracy: 0.7840
Epoch [8/10], Depth: 18, Loss: 0.4824, Accuracy: 0.7820
Epoch [9/10], Depth: 18, Loss: 0.4674, Accuracy: 0.7760
Epoch [10/10], Depth: 18, Loss: 0.4163, Accuracy: 0.7810
Final Depth: 18, Test Accuracy: 0.7810
Epoch [1/10], Depth: 19, Loss: 1.5800, Accuracy: 0.6720
Epoch [2/10], Depth: 19, Loss: 0.9795, Accuracy: 0.7290
Epoch [3/10], Depth: 19, Loss: 0.8651, Accuracy: 0.7270
Epoch [4/10], Depth: 19, Loss: 0.7578, Accuracy: 0.7560
Epoch [5/10], Depth: 19, Loss: 0.7010, Accuracy: 0.7730
Epoch [6/10], Depth: 19, Loss: 0.6043, Accuracy: 0.7790
Epoch [7/10], Depth: 19, Loss: 0.5818, Accuracy: 0.7710
Epoch [8/10], Depth: 19, Loss: 0.5522, Accuracy: 0.7760
Epoch [9/10], Depth: 19, Loss: 0.4883, Accuracy: 0.7940
Epoch [10/10], Depth: 19, Loss: 0.4717, Accuracy: 0.7810
Final Depth: 19, Test Accuracy: 0.7810
Epoch [1/10], Depth: 20, Loss: 1.6040, Accuracy: 0.6650
Epoch [2/10], Depth: 20, Loss: 1.0489, Accuracy: 0.7140
Epoch [3/10], Depth: 20, Loss: 0.9209, Accuracy: 0.7380
Epoch [4/10], Depth: 20, Loss: 0.8057, Accuracy: 0.7740
Epoch [5/10], Depth: 20, Loss: 0.7087, Accuracy: 0.7640
Epoch [6/10], Depth: 20, Loss: 0.6947, Accuracy: 0.7670
Epoch [7/10], Depth: 20, Loss: 0.6226, Accuracy: 0.7680
Epoch [8/10], Depth: 20, Loss: 0.5601, Accuracy: 0.7820
Epoch [9/10], Depth: 20, Loss: 0.5433, Accuracy: 0.7890
Epoch [10/10], Depth: 20, Loss: 0.5273, Accuracy: 0.7820
Final Depth: 20, Test Accuracy: 0.7820
```

Effect of MLP Depth on Test Accuracy



- ✓ This time we vary the sizes of the hidden layers. And experiment with larger and smaller sizes.
- And then we analyze the trade-offs in computational cost and performance of the model.

```
import torch
import torch.nn as nn
import torch.optim as optim
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
import matplotlib.pyplot as plt

class MLPVariableHidden(nn.Module):
    def __init__(self, hidden_size):
        super(MLPVariableHidden, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(50, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, hidden_size),
            nn.BatchNorm1d(hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, 10)
        )

    def forward(self, x):
        return self.model(x)

hidden_sizes = [64, 128, 256, 512, 1024]
accuracies = []

for hidden_size in hidden_sizes:
    model = MLPVariableHidden(hidden_size).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

    num_epochs = 10
    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

        print(f"Epoch [{epoch+1}/{num_epochs}], Hidden Size: {hidden_size}, Loss: {running_loss/len(train_loader):.4f}")

    model.eval()
    correct = 0
    total = 0
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for inputs, labels in test_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    accuracy = correct / total
    accuracies.append(accuracy)
    print(f"Final Hidden Size: {hidden_size}, Test Accuracy: {accuracy:.4f}")


# Plot confusion matrix
```

```
cm = confusion_matrix(all_labels, all_preds)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[str(i) for i in range(10)])
disp.plot(cmap=plt.cm.Blues)
plt.title(f'Confusion Matrix (Hidden Size: {hidden_size})')
plt.show()

plt.figure(figsize=(10, 6))
plt.plot(hidden_sizes, accuracies, marker='o', label='Test Accuracy')
plt.title('Effect of Hidden Layer Size on Test Accuracy')
plt.xlabel('Hidden Layer Size')
plt.ylabel('Accuracy')
plt.grid(True)
plt.legend()
plt.show()
```


Epoch [1/10], Hidden Size: 64, Loss: 1.7655
 Epoch [2/10], Hidden Size: 64, Loss: 1.1582
 Epoch [3/10], Hidden Size: 64, Loss: 0.9146
 Epoch [4/10], Hidden Size: 64, Loss: 0.7936
 Epoch [5/10], Hidden Size: 64, Loss: 0.7111
 Epoch [6/10], Hidden Size: 64, Loss: 0.6553
 Epoch [7/10], Hidden Size: 64, Loss: 0.6075
 Epoch [8/10], Hidden Size: 64, Loss: 0.5684
 Epoch [9/10], Hidden Size: 64, Loss: 0.5492
 Epoch [10/10], Hidden Size: 64, Loss: 0.5329
 Final Hidden Size: 64, Test Accuracy: 0.8090

Confusion Matrix (Hidden Size: 64)



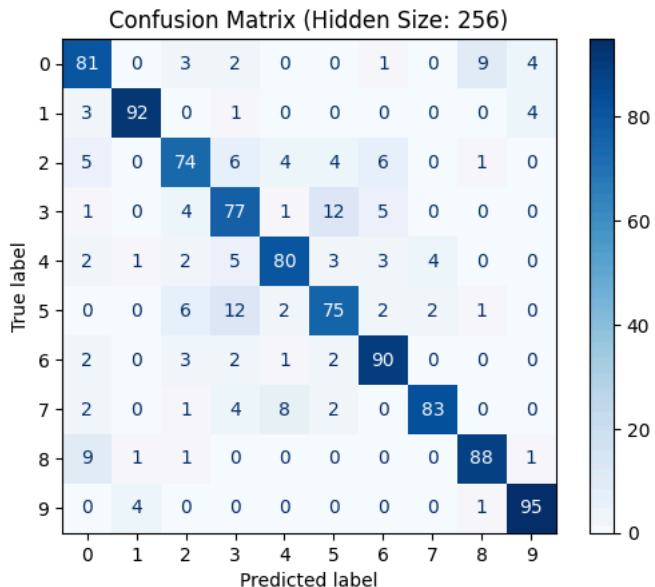
Epoch [1/10], Hidden Size: 128, Loss: 1.6121
 Epoch [2/10], Hidden Size: 128, Loss: 0.9727
 Epoch [3/10], Hidden Size: 128, Loss: 0.7769
 Epoch [4/10], Hidden Size: 128, Loss: 0.6800
 Epoch [5/10], Hidden Size: 128, Loss: 0.6201
 Epoch [6/10], Hidden Size: 128, Loss: 0.5801
 Epoch [7/10], Hidden Size: 128, Loss: 0.5453
 Epoch [8/10], Hidden Size: 128, Loss: 0.5121
 Epoch [9/10], Hidden Size: 128, Loss: 0.4885
 Epoch [10/10], Hidden Size: 128, Loss: 0.4663
 Final Hidden Size: 128, Test Accuracy: 0.8230

Confusion Matrix (Hidden Size: 128)




Epoch [1/10], Hidden Size: 256, Loss: 1.3531
 Epoch [2/10], Hidden Size: 256, Loss: 0.7951
 Epoch [3/10], Hidden Size: 256, Loss: 0.6702
 Epoch [4/10], Hidden Size: 256, Loss: 0.5941
 Epoch [5/10], Hidden Size: 256, Loss: 0.5484
 Epoch [6/10], Hidden Size: 256, Loss: 0.5146
 Epoch [7/10], Hidden Size: 256, Loss: 0.4852
 Epoch [8/10], Hidden Size: 256, Loss: 0.4521
 Epoch [9/10], Hidden Size: 256, Loss: 0.4314
 Epoch [10/10], Hidden Size: 256, Loss: 0.4128

Final Hidden Size: 256, Test Accuracy: 0.8350





Epoch [1/10], Hidden Size: 512, Loss: 1.1041
Epoch [2/10], Hidden Size: 512, Loss: 0.6558
Epoch [3/10], Hidden Size: 512, Loss: 0.5735
Epoch [4/10], Hidden Size: 512, Loss: 0.5151
Epoch [5/10], Hidden Size: 512, Loss: 0.4691
Epoch [6/10], Hidden Size: 512, Loss: 0.4321
Epoch [7/10], Hidden Size: 512, Loss: 0.4069
Epoch [8/10], Hidden Size: 512, Loss: 0.3768
Epoch [9/10], Hidden Size: 512, Loss: 0.3569
Epoch [10/10], Hidden Size: 512, Loss: 0.3398
Final Hidden Size: 512, Test Accuracy: 0.8230



Epoch [1/10], Hidden Size: 1024, Loss: 0.9168
Epoch [2/10], Hidden Size: 1024, Loss: 0.5627
Epoch [3/10], Hidden Size: 1024, Loss: 0.4777
Epoch [4/10], Hidden Size: 1024, Loss: 0.4304
Epoch [5/10], Hidden Size: 1024, Loss: 0.3947
Epoch [6/10], Hidden Size: 1024, Loss: 0.3561
Epoch [7/10], Hidden Size: 1024, Loss: 0.3336
Epoch [8/10], Hidden Size: 1024, Loss: 0.2983
Epoch [9/10], Hidden Size: 1024, Loss: 0.2810
Epoch [10/10], Hidden Size: 1024, Loss: 0.2595
Final Hidden Size: 1024, Test Accuracy: 0.8250





✓ CNN

✓ Define the VGG11 Class

idea taken from: [debugger cafe](#)

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
import torchvision.datasets as datasets
from torch.utils.data import DataLoader

class VGG11(nn.Module):
    def __init__(self):
        super(VGG11, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(),
            nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(256, 512, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(),
            nn.Conv2d(512, 512, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(512, 512, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(),
            nn.Conv2d(512, 512, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
        )

        self._to_linear = nn.Sequential(
            nn.Flatten(),
            nn.Linear(512 * 1 * 1, 4096)
        )

        self.classifier = nn.Sequential(
            nn.Linear(512 * 1 * 1, 4096),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(4096, 10)
        )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1) # Flatten the tensor
        x = self.classifier(x)
```

```
return x
```

▼ Data Loading and Preprocessing

```
import numpy as np

def cifar_loader(batch_size, num_train_per_class=500, num_test_per_class=100):
    normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                      std=[0.229, 0.224, 0.225])

    train_set = datasets.CIFAR10('./data', train=True, download=True,
                                  transform=transforms.Compose([transforms.RandomHorizontalFlip(),
                                                                transforms.RandomCrop(32, 4),
                                                                transforms.ToTensor(), normalize]))

    test_set = datasets.CIFAR10('./data', train=False, download=True,
                                  transform=transforms.Compose([transforms.ToTensor(), normalize]))

    # only include the first 500 images per class for training and 100 for testing
    train_indices = []
    test_indices = []

    for class_label in range(10):
        train_class_indices = np.where(np.array(train_set.targets) == class_label)[0][:num_train_per_class]
        test_class_indices = np.where(np.array(test_set.targets) == class_label)[0][:num_test_per_class]

        train_indices.extend(train_class_indices)
        test_indices.extend(test_class_indices)

    train_subset = torch.utils.data.Subset(train_set, train_indices)
    test_subset = torch.utils.data.Subset(test_set, test_indices)

    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, pin_memory=True)
    test_loader = DataLoader(test_subset, batch_size=batch_size, shuffle=False, pin_memory=True)

    return train_loader, test_loader

batch_size = 64
train_loader, test_loader = cifar_loader(batch_size)
```

📄 Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz
 100%|██████████| 170M/170M [00:05<00:00, 31.1MB/s]
 Extracting ./data/cifar-10-python.tar.gz to ./data
 Files already downloaded and verified

▼ Initialize the model.

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = VGG11().to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
```

▼ Training the model

```
import matplotlib.pyplot as plt

train_losses = []
train_accuracies = []

num_epochs = 30

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0
```

```

for inputs, labels in train_loader:
    inputs, labels = inputs.to(device), labels.to(device)

    # Forward pass
    optimizer.zero_grad()
    outputs = model(inputs)
    loss = criterion(outputs, labels)

    # Backward pass
    loss.backward()
    optimizer.step()

    running_loss += loss.item()

_, predicted = torch.max(outputs, 1)
total += labels.size(0)
correct += (predicted == labels).sum().item()

avg_loss = running_loss / len(train_loader)
accuracy = 100 * correct / total
train_losses.append(avg_loss)
train_accuracies.append(accuracy)

print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}, Accuracy: {accuracy:.2f}%')

```

```

Epoch [1/30], Loss: 2.0166, Accuracy: 24.54%
Epoch [2/30], Loss: 1.7939, Accuracy: 34.16%
Epoch [3/30], Loss: 1.6040, Accuracy: 39.38%
Epoch [4/30], Loss: 1.4678, Accuracy: 46.26%
Epoch [5/30], Loss: 1.3800, Accuracy: 50.10%
Epoch [6/30], Loss: 1.3180, Accuracy: 52.98%
Epoch [7/30], Loss: 1.2758, Accuracy: 55.14%
Epoch [8/30], Loss: 1.2064, Accuracy: 57.48%
Epoch [9/30], Loss: 1.1799, Accuracy: 59.04%
Epoch [10/30], Loss: 1.1493, Accuracy: 59.76%
Epoch [11/30], Loss: 1.0154, Accuracy: 63.98%
Epoch [12/30], Loss: 0.9843, Accuracy: 64.96%
Epoch [13/30], Loss: 0.9507, Accuracy: 66.84%
Epoch [14/30], Loss: 0.9103, Accuracy: 68.16%
Epoch [15/30], Loss: 0.8894, Accuracy: 69.92%
Epoch [16/30], Loss: 0.8477, Accuracy: 70.50%
Epoch [17/30], Loss: 0.7918, Accuracy: 73.24%
Epoch [18/30], Loss: 0.8004, Accuracy: 73.38%
Epoch [19/30], Loss: 0.7673, Accuracy: 73.24%
Epoch [20/30], Loss: 0.7453, Accuracy: 74.86%
Epoch [21/30], Loss: 0.7467, Accuracy: 74.86%
Epoch [22/30], Loss: 0.6601, Accuracy: 77.96%
Epoch [23/30], Loss: 0.6852, Accuracy: 77.26%
Epoch [24/30], Loss: 0.6190, Accuracy: 78.86%
Epoch [25/30], Loss: 0.5778, Accuracy: 80.08%
Epoch [26/30], Loss: 0.5913, Accuracy: 80.20%
Epoch [27/30], Loss: 0.6047, Accuracy: 79.58%
Epoch [28/30], Loss: 0.5404, Accuracy: 81.78%
Epoch [29/30], Loss: 0.5140, Accuracy: 83.20%
Epoch [30/30], Loss: 0.5256, Accuracy: 81.72%

```

✓ Save the model:

code suggested by Google Gemini.

```

torch.save(model, 'vgg11_full_model.pth')
print("Full model saved to vgg11_full_model.pth")
from google.colab import files
files.download('vgg11_full_model.pth')

```

```

Full model saved to vgg11_full_model.pth

```

✓ Training loss and accuracy:

```

epochs = range(1, num_epochs + 1)

plt.figure(figsize=(12, 6))

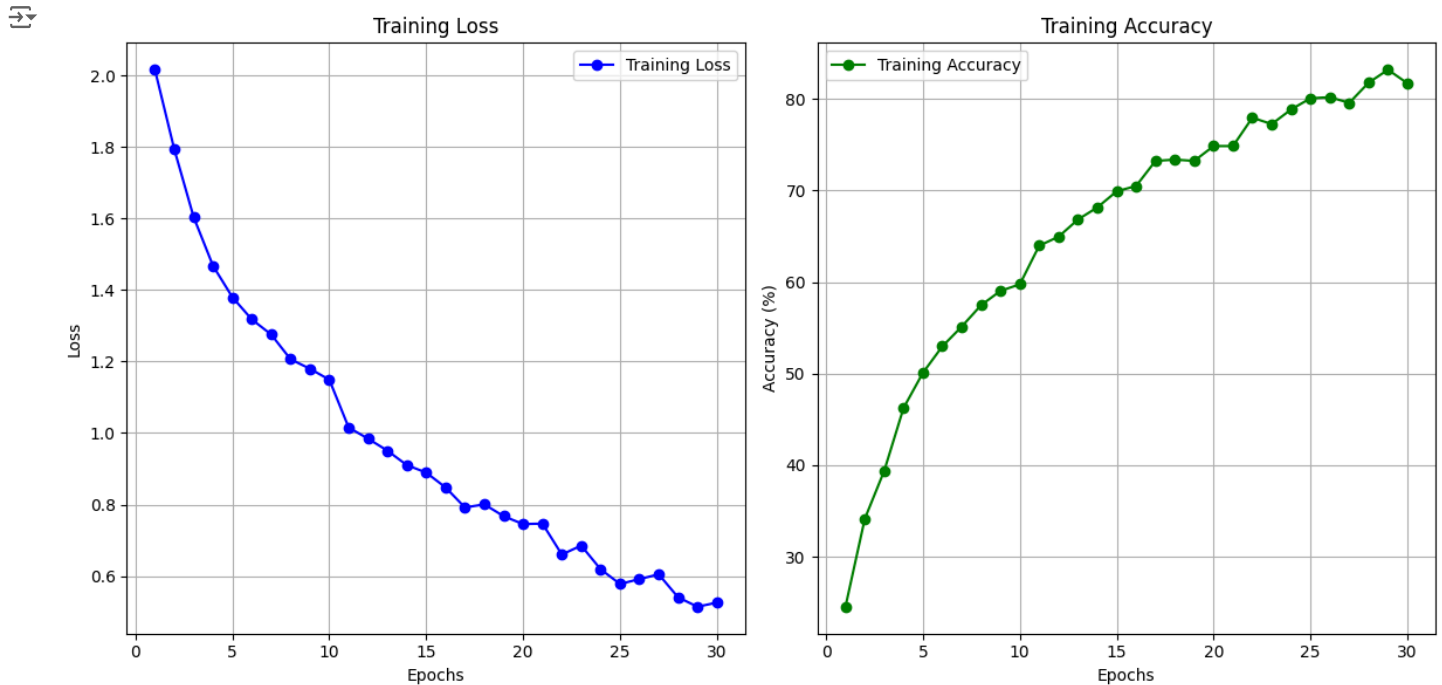
```



```
plt.subplot(1, 2, 1)
plt.plot(epochs, train_losses, marker='o', color='b', label='Training Loss')
plt.title('Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid(True)
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, train_accuracies, marker='o', color='g', label='Training Accuracy')
plt.title('Training Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy (%)')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()
```



✓ Evaluate the model:

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

model.eval()
correct = 0
total = 0
all_preds = []
all_labels = []

with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())


accuracy = correct / total
print(f"Test Accuracy: {accuracy:.4f}")
```

```
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
cm = confusion_matrix(all_labels, all_preds)
cm_display = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
```

```
cm_display.plot(cmap='Blues')
plt.title('Confusion Matrix for VGG11')
plt.show()
```

Test Accuracy: 0.7040



✓ Experiment with the model by adding and removing layers.

we will define a custome VG11 class for this.

```
# modified VGG architectures (idea from : https://debuggercafe.com/implementing-vgg11-from-scratch-using-pytorch/)
```

```
class ModifiedVGG(nn.Module):
    def __init__(self, layers_config):
        super(ModifiedVGG, self).__init__()
        layers = []
        in_channels = 3
        for config in layers_config:
            if config == 'M':
                layers += [nn.MaxPool2d(kernel_size=2, stride=2)]
            else:
                layers += [
                    nn.Conv2d(in_channels, config, kernel_size=3, padding=1),
                    nn.BatchNorm2d(config),
                    nn.ReLU(inplace=True)
                ]
                in_channels = config
        self.features = nn.Sequential(*layers)

    def compute_fc_input_size(x):
        x = self.features(x)
        return x.view(x.size(0), -1)

dummy_input = torch.zeros(1, 3, 32, 32)
self.fc_input_size = compute_fc_input_size(dummy_input).size(1)

self.classifier = nn.Sequential(
    nn.Linear(self.fc_input_size, 4096),
    nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(4096, 4096),
    nn.ReLU(),
```

```

        nn.Dropout(0.5),
        nn.Linear(4096, 10)
    )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x

layer_configs = {
    "VGG_Light": [64, 'M', 128, 'M', 256, 'M', 512, 'M'],
    "VGG_Full": [64, 'M', 128, 'M', 256, 256, 'M', 512, 512, 'M', 512, 512, 'M'],
    "VGG_Deep": [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 'M', 512, 512, 512, 'M']
}

results = {}
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

batch_size = 64
train_loader, test_loader = cifar_loader(batch_size)

for name, config in layer_configs.items():
    print(f"\nTraining {name}...")
    model = ModifiedVGG(config).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

    num_epochs = 30
    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()
        print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(train_loader):.4f}")

    model.eval()
    all_preds = []
    all_labels = []
    with torch.no_grad():
        for inputs, labels in test_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted = torch.max(outputs, 1)
            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    acc = sum(1 for x, y in zip(all_preds, all_labels) if x == y) / len(all_labels)
    cm = confusion_matrix(all_labels, all_preds)
    results[name] = {'accuracy': acc, 'confusion_matrix': cm}
    print(f"{name} Accuracy: {acc:.4f}")

fig, axes = plt.subplots(1, 3, figsize=(18, 6))
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for ax, (name, result) in zip(axes, results.items()):
    cm_display = ConfusionMatrixDisplay(confusion_matrix=result['confusion_matrix'], display_labels=class_names)
    cm_display.plot(ax=ax, cmap='Blues', colorbar=False)
    ax.set_title(f"{name} Confusion Matrix")
plt.tight_layout()
plt.show()

results

```

Files already downloaded and verified
Files already downloaded and verified

```
Training VGG_Light...
Epoch [1/30], Loss: 2.0073
Epoch [2/30], Loss: 1.7059
Epoch [3/30], Loss: 1.6278
Epoch [4/30], Loss: 1.5147
Epoch [5/30], Loss: 1.4475
Epoch [6/30], Loss: 1.3815
Epoch [7/30], Loss: 1.3442
Epoch [8/30], Loss: 1.2764
Epoch [9/30], Loss: 1.2258
Epoch [10/30], Loss: 1.2398
Epoch [11/30], Loss: 1.1812
Epoch [12/30], Loss: 1.1641
Epoch [13/30], Loss: 1.0954
Epoch [14/30], Loss: 1.0569
Epoch [15/30], Loss: 1.0449
Epoch [16/30], Loss: 0.9909
Epoch [17/30], Loss: 0.9842
Epoch [18/30], Loss: 0.9685
Epoch [19/30], Loss: 0.9411
Epoch [20/30], Loss: 0.9059
Epoch [21/30], Loss: 0.8828
Epoch [22/30], Loss: 0.8675
Epoch [23/30], Loss: 0.8615
Epoch [24/30], Loss: 0.8204
Epoch [25/30], Loss: 0.7983
Epoch [26/30], Loss: 0.7795
Epoch [27/30], Loss: 0.7742
Epoch [28/30], Loss: 0.7298
Epoch [29/30], Loss: 0.7288
Epoch [30/30], Loss: 0.7435
VGG_Light Accuracy: 0.6660
```




```
Training VGG_Full...
Epoch [1/30], Loss: 2.0195
Epoch [2/30], Loss: 1.7245
Epoch [3/30], Loss: 1.5631
Epoch [4/30], Loss: 1.4857
Epoch [5/30], Loss: 1.3918
Epoch [6/30], Loss: 1.3131
Epoch [7/30], Loss: 1.2621
Epoch [8/30], Loss: 1.1862
Epoch [9/30], Loss: 1.1328
Epoch [10/30], Loss: 1.0784
Epoch [11/30], Loss: 1.0115
Epoch [12/30], Loss: 0.9916
Epoch [13/30], Loss: 0.9212
Epoch [14/30], Loss: 0.8869
Epoch [15/30], Loss: 0.8839
Epoch [16/30], Loss: 0.8476
Epoch [17/30], Loss: 0.7727
Epoch [18/30], Loss: 0.7865
Epoch [19/30], Loss: 0.7636
Epoch [20/30], Loss: 0.7167
Epoch [21/30], Loss: 0.7518
Epoch [22/30], Loss: 0.6855
Epoch [23/30], Loss: 0.6588
Epoch [24/30], Loss: 0.6342
Epoch [25/30], Loss: 0.6024
Epoch [26/30], Loss: 0.5678
Epoch [27/30], Loss: 0.5377
Epoch [28/30], Loss: 0.5390
Epoch [29/30], Loss: 0.5025
Epoch [30/30], Loss: 0.5374
VGG_Full Accuracy: 0.6870
```

```
Training VGG_Deep...
Epoch [1/30], Loss: 2.0425
Epoch [2/30], Loss: 1.7501
Epoch [3/30], Loss: 1.6213
Epoch [4/30], Loss: 1.5043
Epoch [5/30], Loss: 1.4728
Epoch [6/30], Loss: 1.3740
Epoch [7/30], Loss: 1.3240
Epoch [8/30], Loss: 1.2165
Epoch [9/30], Loss: 1.1592
Epoch [10/30], Loss: 1.0765
Epoch [11/30], Loss: 1.1188
Epoch [12/30], Loss: 0.9757
Epoch [13/30], Loss: 0.9426
Epoch [14/30], Loss: 0.9410
```

```

Epoch [15/30], Loss: 0.8962
Epoch [16/30], Loss: 0.8201
Epoch [17/30], Loss: 0.7881
Epoch [18/30], Loss: 0.7345
Epoch [19/30], Loss: 0.7307
Epoch [20/30], Loss: 0.7081
Epoch [21/30], Loss: 0.7032
Epoch [22/30], Loss: 0.6438
Epoch [23/30], Loss: 0.6428
Epoch [24/30], Loss: 0.6424
Epoch [25/30], Loss: 0.5928
Epoch [26/30], Loss: 0.5356
Epoch [27/30], Loss: 0.6016
Epoch [28/30], Loss: 0.4859
Epoch [29/30], Loss: 0.5100
Epoch [30/30], Loss: 0.4892
VGG_Deep Accuracy: 0.6870

```



```

{'VGG_Light': {'accuracy': 0.666,
'confusion_matrix': array([[62, 13, 0, 2, 1, 0, 3, 1, 12, 6],
[ 1, 95, 0, 0, 0, 0, 0, 0, 1, 1, 2],
[20, 5, 30, 14, 2, 14, 9, 2, 2, 2],
[ 5, 8, 1, 45, 5, 20, 7, 2, 4, 3],
[ 3, 3, 3, 13, 52, 4, 5, 11, 6, 0],
[ 0, 5, 2, 20, 2, 66, 1, 1, 3, 0],
[ 1, 6, 3, 8, 3, 2, 75, 1, 0, 1],
[ 1, 5, 0, 2, 1, 7, 1, 79, 2, 2],
[ 6, 9, 0, 0, 0, 0, 0, 0, 82, 3],
[ 0, 13, 0, 0, 0, 1, 0, 0, 6, 80]])},
'VGG_Full': {'accuracy': 0.687,
'confusion_matrix': array([[69, 2, 2, 2, 1, 0, 2, 0, 13, 9],
[ 4, 64, 0, 0, 0, 1, 0, 0, 6, 25],
[ 9, 1, 49, 13, 11, 7, 1, 3, 4, 2],
[ 8, 0, 6, 59, 5, 9, 2, 2, 2, 7],
[ 4, 1, 4, 7, 70, 1, 0, 9, 2, 2],
[ 0, 0, 8, 31, 2, 48, 2, 2, 1, 6],
[ 0, 0, 3, 13, 9, 2, 68, 1, 1, 3],
[ 2, 0, 0, 6, 4, 3, 0, 78, 3, 4],
[ 4, 1, 0, 0, 2, 0, 0, 0, 89, 4],
[ 1, 1, 0, 0, 0, 0, 1, 0, 4, 93]])},
'VGG_Deep': {'accuracy': 0.687,
'confusion_matrix': array([[75, 3, 3, 2, 0, 0, 0, 0, 8, 9],
[ 1, 91, 0, 0, 0, 0, 0, 0, 2, 6],
[15, 0, 67, 8, 5, 1, 1, 0, 1, 2],
[10, 1, 5, 62, 8, 2, 3, 3, 1, 5],
[ 6, 1, 10, 12, 68, 2, 0, 0, 1, 0],
[ 3, 0, 13, 34, 4, 39, 1, 3, 1, 2],
[ 4, 2, 20, 14, 1, 0, 58, 1, 0, 0],
[ 6, 2, 7, 9, 14, 1, 0, 55, 0, 6],
[ 9, 2, 0, 0, 0, 0, 0, 0, 86, 3],
[ 1, 8, 0, 1, 0, 0, 0, 0, 4, 86]])}

```

Double-click (or enter) to edit

✓ Modified VGG Class with Adjustable Kernel Sizes:

```
class ModifiedVGGWithKernels(nn.Module):
    def __init__(self, layers_config, kernel_size=3):
        super(ModifiedVGGWithKernels, self).__init__()
        layers = []
        in_channels = 3
        for config in layers_config:
            if config == 'M':
                layers += [nn.MaxPool2d(kernel_size=2, stride=2)]
            else:
                layers += [
                    nn.Conv2d(in_channels, config, kernel_size=kernel_size, padding=kernel_size//2),
                    nn.BatchNorm2d(config),
                    nn.ReLU(inplace=True)
                ]
                in_channels = config
        self.features = nn.Sequential(*layers)

    def compute_fc_input_size(x):
        x = self.features(x)
        return x.view(x.size(0), -1)

    dummy_input = torch.zeros(1, 3, 32, 32)
    self.fc_input_size = compute_fc_input_size(dummy_input).size(1)

    self.classifier = nn.Sequential(
        nn.Linear(self.fc_input_size, 4096),
        nn.ReLU(),
        nn.Dropout(0.5),
        nn.Linear(4096, 4096),
        nn.ReLU(),
        nn.Dropout(0.5),
        nn.Linear(4096, 10)
    )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x
```

✓ Experiment with different Kernel sizes:

```
kernel_sizes = [2, 3, 5, 7]
results_kernel_sizes = {}

for kernel_size in kernel_sizes:
    print(f"\nTraining with kernel size {kernel_size}x{kernel_size}...")

    model = ModifiedVGGWithKernels([64, 'M', 128, 'M', 256, 'M', 512, 'M'], kernel_size).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

    num_epochs = 30
    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()
```

```

print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(train_loader):.4f}")

model.eval()
all_preds = []
all_labels = []
with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, predicted = torch.max(outputs, 1)
        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

acc = sum(1 for x, y in zip(all_preds, all_labels) if x == y) / len(all_labels)
cm = confusion_matrix(all_labels, all_preds)
results_kernel_sizes[kernel_size] = {'accuracy': acc, 'confusion_matrix': cm}
print(f"Kernel size {kernel_size}x{kernel_size} Accuracy: {acc:.4f}")

fig, axes = plt.subplots(1, len(kernel_sizes), figsize=(18, 6))
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for ax, (kernel_size, result) in zip(axes, results_kernel_sizes.items()):
    cm_display = ConfusionMatrixDisplay(confusion_matrix=result['confusion_matrix'], display_labels=class_names)
    cm_display.plot(ax=ax, cmap='Blues', colorbar=False)
    ax.set_title(f"Kernel {kernel_size}x{kernel_size}")
plt.tight_layout()
plt.show()

results_kernel_sizes

```